

In [5]: ▶

```
1 import numpy as np
2
3 from numpy import array
4 from keras.models import Sequential
5 from keras.layers import LSTM
6 from keras.layers import Dense
```

In [6]: ▶

```
1 from sklearn.datasets import fetch_california_housing
```

In [6]: ▶

```
1 ca_house_db=fetch_california_housing()
2 print (ca_house_db.data.shape)
3 print (ca_house_db.target.shape)
4 print (ca_house_db.feature_names)
5 print (ca_house_db.DESCR)
6 print (ca_house_db)
```

```
(20640, 8)
(20640,)
['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup',
'Latitude', 'Longitude']
.. _california_housing_dataset:
```

California Housing dataset

****Data Set Characteristics:****

:Number of Instances: 20640

:Number of Attributes: 8 numeric, predictive attributes and the target

:Attribute Information:

- MedInc median income in block group
- HouseAge median house age in block group
- AveRooms average number of rooms per household
- AveBedrms average number of bedrooms per household
- Population block group population
- AveOccup average number of household members
- Latitude block group latitude
- Longitude block group longitude

:Missing Attribute Values: None

This dataset was obtained from the StatLib repository.

https://www.dcc.fc.up.pt/~ltorgo/Regression/cal_housing.html (https://www.dcc.fc.up.pt/~ltorgo/Regression/cal_housing.html)

The target variable is the median house value for California districts, expressed in hundreds of thousands of dollars (\$100,000).

This dataset was derived from the 1990 U.S. census, using one row per census

block group. A block group is the smallest geographical unit for which the U.S.

Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people).

An household is a group of people residing within a home. Since the average

number of rooms and bedrooms in this dataset are provided per household, these

columns may take surprisingly large values for block groups with few households

and many empty houses, such as vacation resorts.

It can be downloaded/loaded using the

:func:`sklearn.datasets.fetch_california_housing` function.

.. topic:: References

- Pace, R. Kelley and Ronald Barry, Sparse Spatial Autoregressions,

Statistics and Probability Letters, 33 (1997) 291-297

```
{'data': array([[ 8.3252, 41., 6.98412698, ..., 2.55555556,
                 37.88, -122.23],
                [ 8.3014, 21., 6.23813708, ..., 2.10984183,
                 37.86, -122.22],
                [ 7.2574, 52., 8.28813559, ..., 2.80225989,
                 37.85, -122.24],
                ...,
                [ 1.7, 17., 5.20554273, ..., 2.3256351,
                 39.43, -121.22],
                [ 1.8672, 18., 5.32951289, ..., 2.12320917,
                 39.43, -121.32],
                [ 2.3886, 16., 5.25471698, ..., 2.61698113,
                 39.37, -121.24]])}, 'target': array([4.526, 3.585,
3.521, ..., 0.923, 0.847, 0.894]), 'frame': None, 'target_names': ['MedHouseVal'], 'feature_names': ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude'], 'DESCR': '.. _california_housing_dataset:\n\nCalifornia Housing dataset\n-----\n\n**Data Set Characteristics:**\n\n: Number of Instances: 20640\n\n: Number of Attributes: 8 numeric, predictive attributes and the target\n\n: Attribute Information:\n\n- MedInc median income in block group\n- HouseAge median house age in block group\n- AveRooms average number of rooms per household\n- AveBedrms average number of bedrooms per household\n- Population block group population\n- AveOccup average number of household members\n- Latitude block group latitude\n- Longitude block group longitude\n\n: Missing Attribute Values: None\n\nThis dataset was obtained from the StatLib repository.\nhttps://www.dcc.fc.up.pt/~ltorgo/Regression/cal_housing.html\n\nThe target variable is the median house value for California districts, expressed in hundreds of thousands of dollars ($100,000).\n\nThis dataset was derived from the 1990 U.S. census, using one row per census block group. A block group is the smallest geographical unit for which the U.S. Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people).\n\nAn household is a group of people residing within a home. Since the average number of rooms and bedrooms in this dataset are provided per household, these columns may take surprisingly large values for block groups with few households and many empty houses, such as vacation resorts.\n\nIt can be downloaded/loaded using the\n:func:`sklearn.datasets.fetch_california_housing` function.\n\n.. topic:: References\n\n- Pace, R. Kelley and Ronald Barry, Sparse Spatial Autoregressions, Statistics and Probability Letters, 33 (1997) 291-297\n'}
```

```
In [10]: 1 housing = fetch_california_housing()
2 X = housing.data
3 y = housing.target
```

```
In [15]: 1 from sklearn.model_selection import train_test_split
```

In [16]: ▶

```
1 # Split the data into training (50%) and validation (50%) sets
2 X_train_val, X_test, y_train_val, y_test = train_test_split(housing.data
3
4 # Split the training set into training (80%) and validation (20%) subsets
5 X_train, X_val, y_train, y_val = train_test_split(X_train_val, y_train_val,
6
7 # Print the shapes of the resulting subsets
8 print("Training set: X_train = {}, y_train = {}".format(X_train.shape,
9 print("Validation set: X_val = {}, y_val = {}".format(X_val.shape, y_val
10 print("Testing set: X_test = {}, y_test = {}".format(X_test.shape, y_test
11
```

Training set: X_train = (8256, 8), y_train = (8256,)
Validation set: X_val = (2064, 8), y_val = (2064,)
Testing set: X_test = (10320, 8), y_test = (10320,)

```

In [9]: 1 import tensorflow as tf
2 from sklearn.datasets import fetch_california_housing
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import r2_score
5
6 # Load data
7 housing = fetch_california_housing()
8 X_train_full, X_test, y_train_full, y_test = train_test_split(housing.
9 X_train, X_valid, y_train, y_valid = train_test_split(X_train_full, y_
10
11 # Define model
12 model = tf.keras.models.Sequential([
13     tf.keras.layers.Dense(64, activation="relu", input_shape=X_train.s
14     tf.keras.layers.Dense(64, activation="relu"),
15     tf.keras.layers.Dense(64, activation="relu"),
16     tf.keras.layers.Dense(1, activation="linear")
17 ])
18
19 # Compile model
20 model.compile(loss="mse", optimizer="adam")
21
22 # Train model
23 history = model.fit(X_train, y_train, epochs=100, validation_data=(X_v
24
25 # Predict on validation set
26 y_valid_pred = model.predict(X_valid)
27
28 # Calculate R2 score on validation set
29 r2 = r2_score(y_valid, y_valid_pred)
30 print(f"Set 1 R2 score: {r2}")
31

```

```

Epoch 1/100
363/363 [=====] - 2s 3ms/step - loss: 46.6198
- val_loss: 5.5429
Epoch 2/100
363/363 [=====] - 1s 2ms/step - loss: 1.3149
- val_loss: 4.4744
Epoch 3/100
363/363 [=====] - 1s 2ms/step - loss: 4.5052
- val_loss: 5.6807
Epoch 4/100
363/363 [=====] - 1s 3ms/step - loss: 73.1594
- val_loss: 2.7875
Epoch 5/100
363/363 [=====] - 1s 2ms/step - loss: 1.0399
- val_loss: 3.0086
Epoch 6/100
363/363 [=====] - 1s 3ms/step - loss: 1.1106
- val_loss: 3.3342
Epoch 7/100
363/363 [=====] - 1s 3ms/step - loss: 1.1106
- val_loss: 3.3342

```

```

In [1]: 1 import tensorflow as tf
2 from sklearn.datasets import fetch_california_housing
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import r2_score
5
6 # Load data
7 housing = fetch_california_housing()
8 X_train_full, X_test, y_train_full, y_test = train_test_split(housing.
9 X_train, X_valid, y_train, y_valid = train_test_split(X_train_full, y
10
11 # Define model
12 model = tf.keras.models.Sequential([
13     tf.keras.layers.Dense(128, activation="sigmoid", input_shape=X_tra
14     tf.keras.layers.Dense(128, activation="sigmoid"),
15     tf.keras.layers.Dense(1, activation="linear")
16 ])
17
18 # Compile model
19 model.compile(loss="mse", optimizer="sgd")
20
21 # Train model
22 history = model.fit(X_train, y_train, epochs=100, validation_data=(X_v
23
24 # Predict on validation set
25 y_valid_pred = model.predict(X_valid)
26
27 # Calculate R2 score on validation set
28 r2 = r2_score(y_valid, y_valid_pred)
29 print(f"Set 2 R2 score: {r2}")
30

```

Epoch 62/100

363/363 [=====] - 1s 3ms/step - loss: 1.3447
- val_loss: 1.3416

Epoch 63/100

363/363 [=====] - 1s 3ms/step - loss: 1.3427
- val_loss: 1.3253

Epoch 64/100

363/363 [=====] - 1s 3ms/step - loss: 1.3435
- val_loss: 1.3147

Epoch 65/100

363/363 [=====] - 1s 2ms/step - loss: 1.3438
- val_loss: 1.3204

Epoch 66/100

363/363 [=====] - 1s 2ms/step - loss: 1.3434
- val_loss: 1.3290

Epoch 67/100

363/363 [=====] - 1s 2ms/step - loss: 1.3430
- val_loss: 1.3145

Epoch 68/100

363/363 [=====] - 1s 3ms/step - loss: 1.3434

```
In [10]: 1 from sklearn.metrics import r2_score
2
3 # Define the model
4 model = tf.keras.models.Sequential([
5     tf.keras.layers.Dense(64, activation='relu', input_shape=(8,)),
6     tf.keras.layers.Dense(64, activation='relu'),
7     tf.keras.layers.Dense(64, activation='relu'),
8     tf.keras.layers.Dense(1)
9 ])
10
11 # Compile the model
12 model.compile(optimizer='adam', loss='mse', metrics=['mae'])
13
14 # Train the model on the entire training set
15 model.fit(X_train, y_train, epochs=50, batch_size=128, verbose=0)
16
17 # Evaluate the model on the testing set
18 y_pred = model.predict(X_test)
19 r2 = r2_score(y_test, y_pred)
20 print("R2 score on testing set:", r2)
21
```

162/162 [=====] - 0s 1ms/step
R2 score on testing set: 0.4053987785465406

```
In [11]: 1 from sklearn.metrics import r2_score
2
3 # Define the model
4 model = tf.keras.models.Sequential([
5     tf.keras.layers.Dense(64, activation='relu', input_shape=(8,)),
6     tf.keras.layers.Dense(64, activation='relu'),
7     tf.keras.layers.Dense(64, activation='relu'),
8     tf.keras.layers.Dense(1)
9 ])
10
11 # Compile the model
12 model.compile(optimizer='adam', loss='mse', metrics=['mae'])
13
14 # Train the model on the entire training set
15 model.fit(X_train, y_train, epochs=50, batch_size=128, verbose=0)
16
17 # Evaluate the model on the testing set
18 y_pred = model.predict(X_test)
19 r2 = r2_score(y_test, y_pred)
20 print("R2 score on testing set:", r2)
21
```

162/162 [=====] - 0s 1ms/step
R2 score on testing set: 0.5257819877370892


```
In [12]: ▶ 1 # Apply top-ranked model over testing samples
          2 y_pred_test = model.predict(X_test)
          3
          4 # Calculate absolute errors
          5 abs_errors = np.abs(y_test - y_pred_test).flatten()
          6
          7 # Sort errors in descending order and select top 10
          8 largest_errors_idx = np.argsort(abs_errors)[:,-1][:10]
          9 largest_errors = abs_errors[largest_errors_idx]
         10

162/162 [=====] - 0s 1ms/step
```

The R2 score on the testing set for the selected model is 0.525, which indicates a good fit of the model on the testing data.